

Machine Learning

Abdelhak Mahmoudi
abdelhak.mahmoudi@um5.ac.ma

INPT - 2020

Content

1. The Big Picture

2. Supervised Learning

- Linear Regression, Logistic Regression, Support Vector Machines, Trees, Random Forests, Boosting, Artificial Neural Networks

3. Unsupervised Learning

- Principal Component Analysis, K-means, Mean Shift

The Big Picture

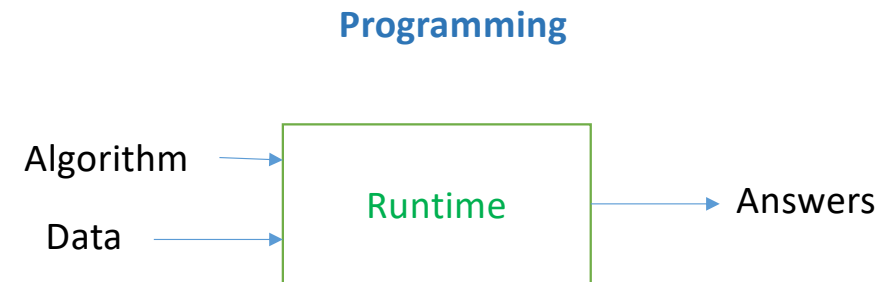
- The Big Picture of ML !
- Terminologies
- How can I Apply?
- How can I Learn?

The Big Picture!

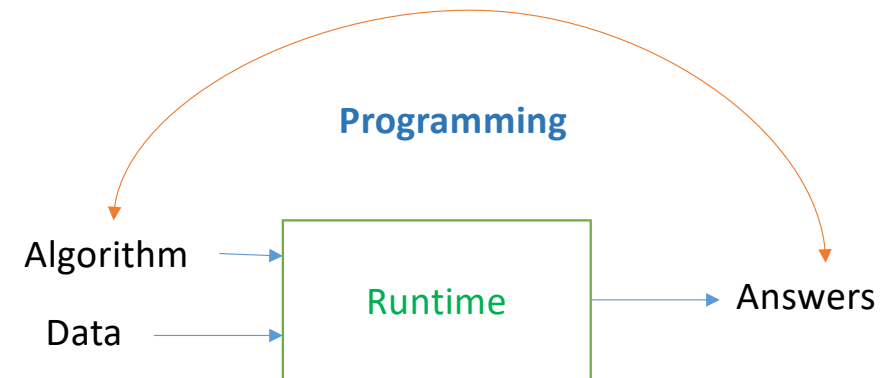
Forbes: “The Top 10 AI And Machine Learning Use Cases Everyone Should Know About”

1. Data Security,
2. Personal Security,
3. Financial Trading,
4. Healthcare,
5. Marketing personalization,
6. Fraud Detection,
7. Recommendations,
8. Online Search,
9. Natural Language Processing (NLP),
10. Smart Cars


The Big Picture!



The Big Picture!

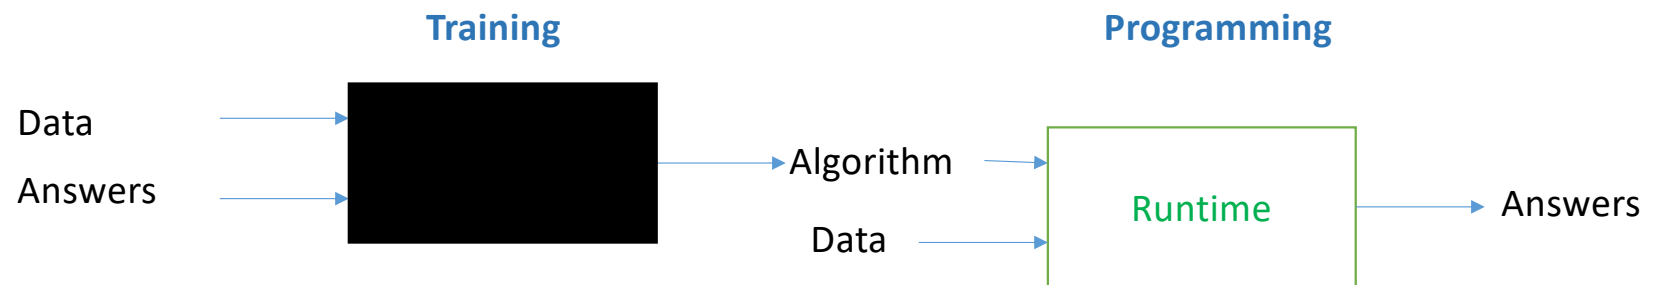


The Big Picture!




The Big Picture!

Machine Learning




The Big Picture!

Machine Learning



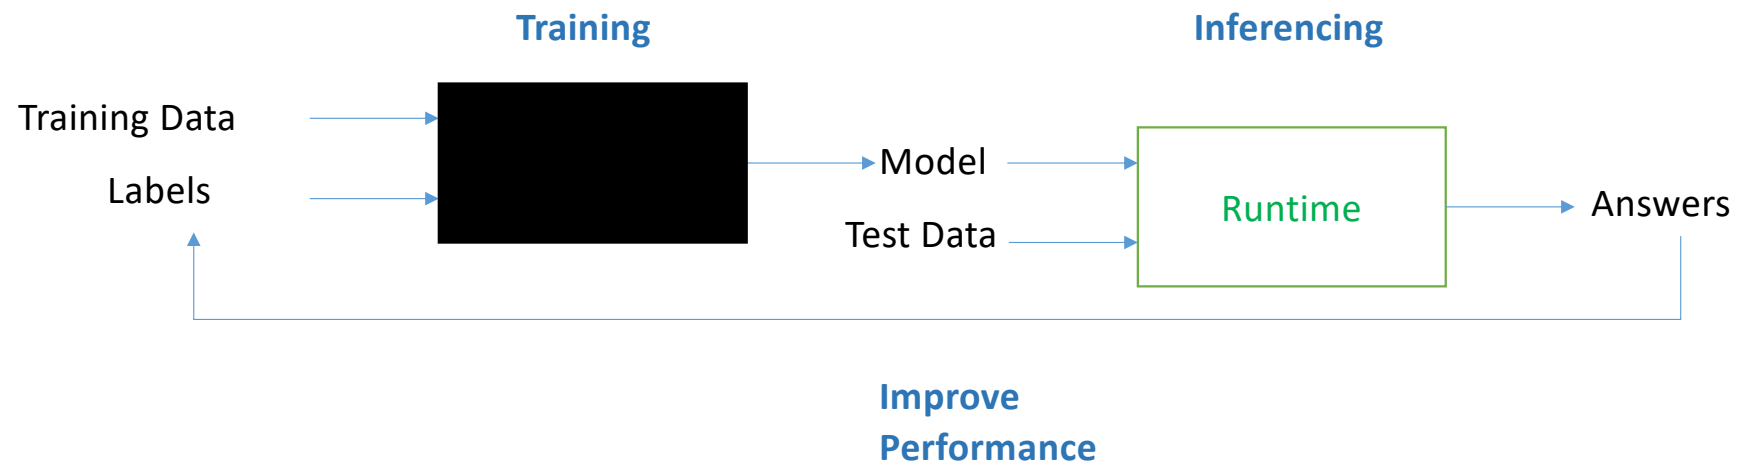
The Big Picture!

Machine Learning



The Big Picture!


Machine Learning



The Big Picture!

- **Data**

- Example $x^{(i)}$
 - Row/Instance/Input/Observation/Record/Point/Sample/Entity
- Feature $x^{(i)}_j$
 - Columns/Variable/Predictor/Characteristic/Field/Attribute
 - Quantitative (numeric, continue)
 - Qualitative (textual, category)
- Dimension, Visualization
 - m Examples: $i = 1..m$
 - n Features: $j = 1..n$
- Output : $y_i = x^{(i)}_k$ (k in $1..n$)
 - target/class/output
 - For each example (0/1)



The Big Picture!


- Data

- Structured

- CSV, XML, JSON, XLSX, etc.

- Unstructured

- DOC, HTML, PDF, PNG, MP3, MP4, etc.

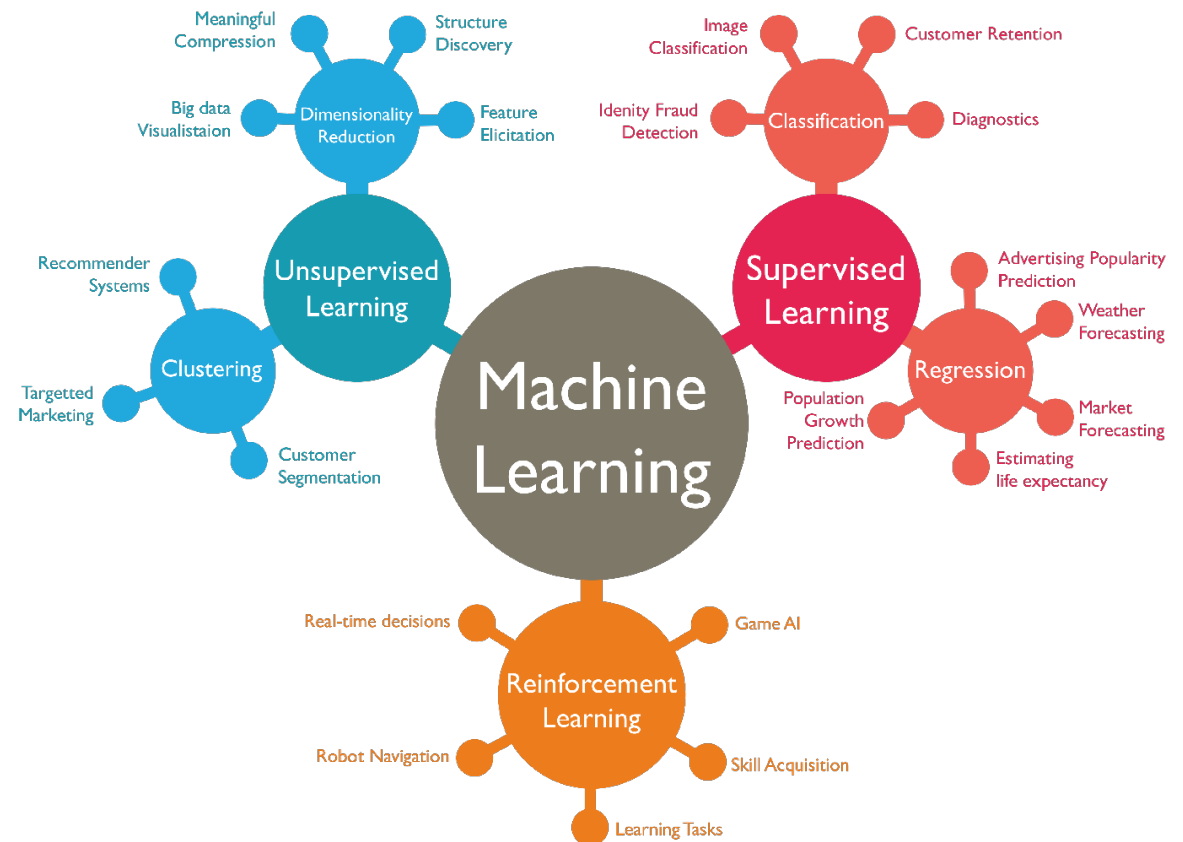


Text, Image, son

The Big Picture!

- **Types of Learning**


- Supervised
 - Classification
 - Regression
- Unsupervised
 - Dimensionality Reduction
 - Clustering
- Semi-supervised
 - Little supervised data
- Reinforcement




The Big Picture!

Supervised Learning

Classification (y is discrete)



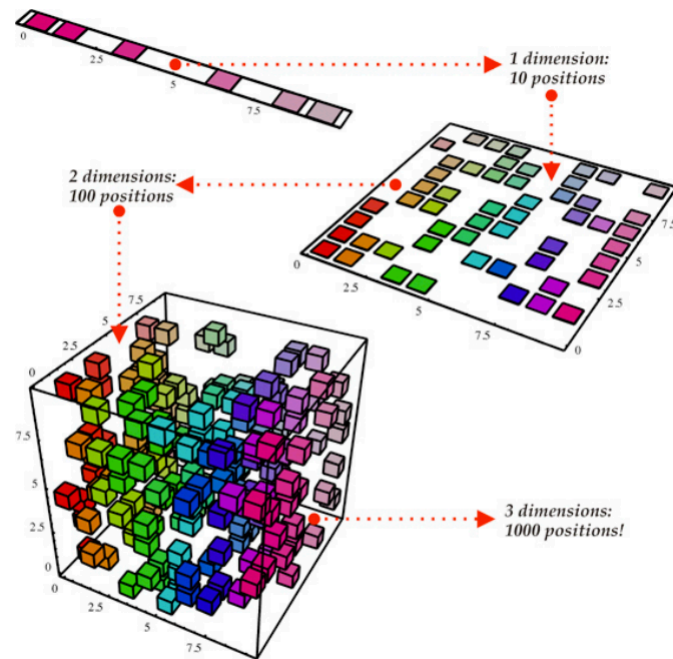
Regression (y is continuous)




The Big Picture!

Unsupervised Learning (y absent)



Dimensionality Reduction



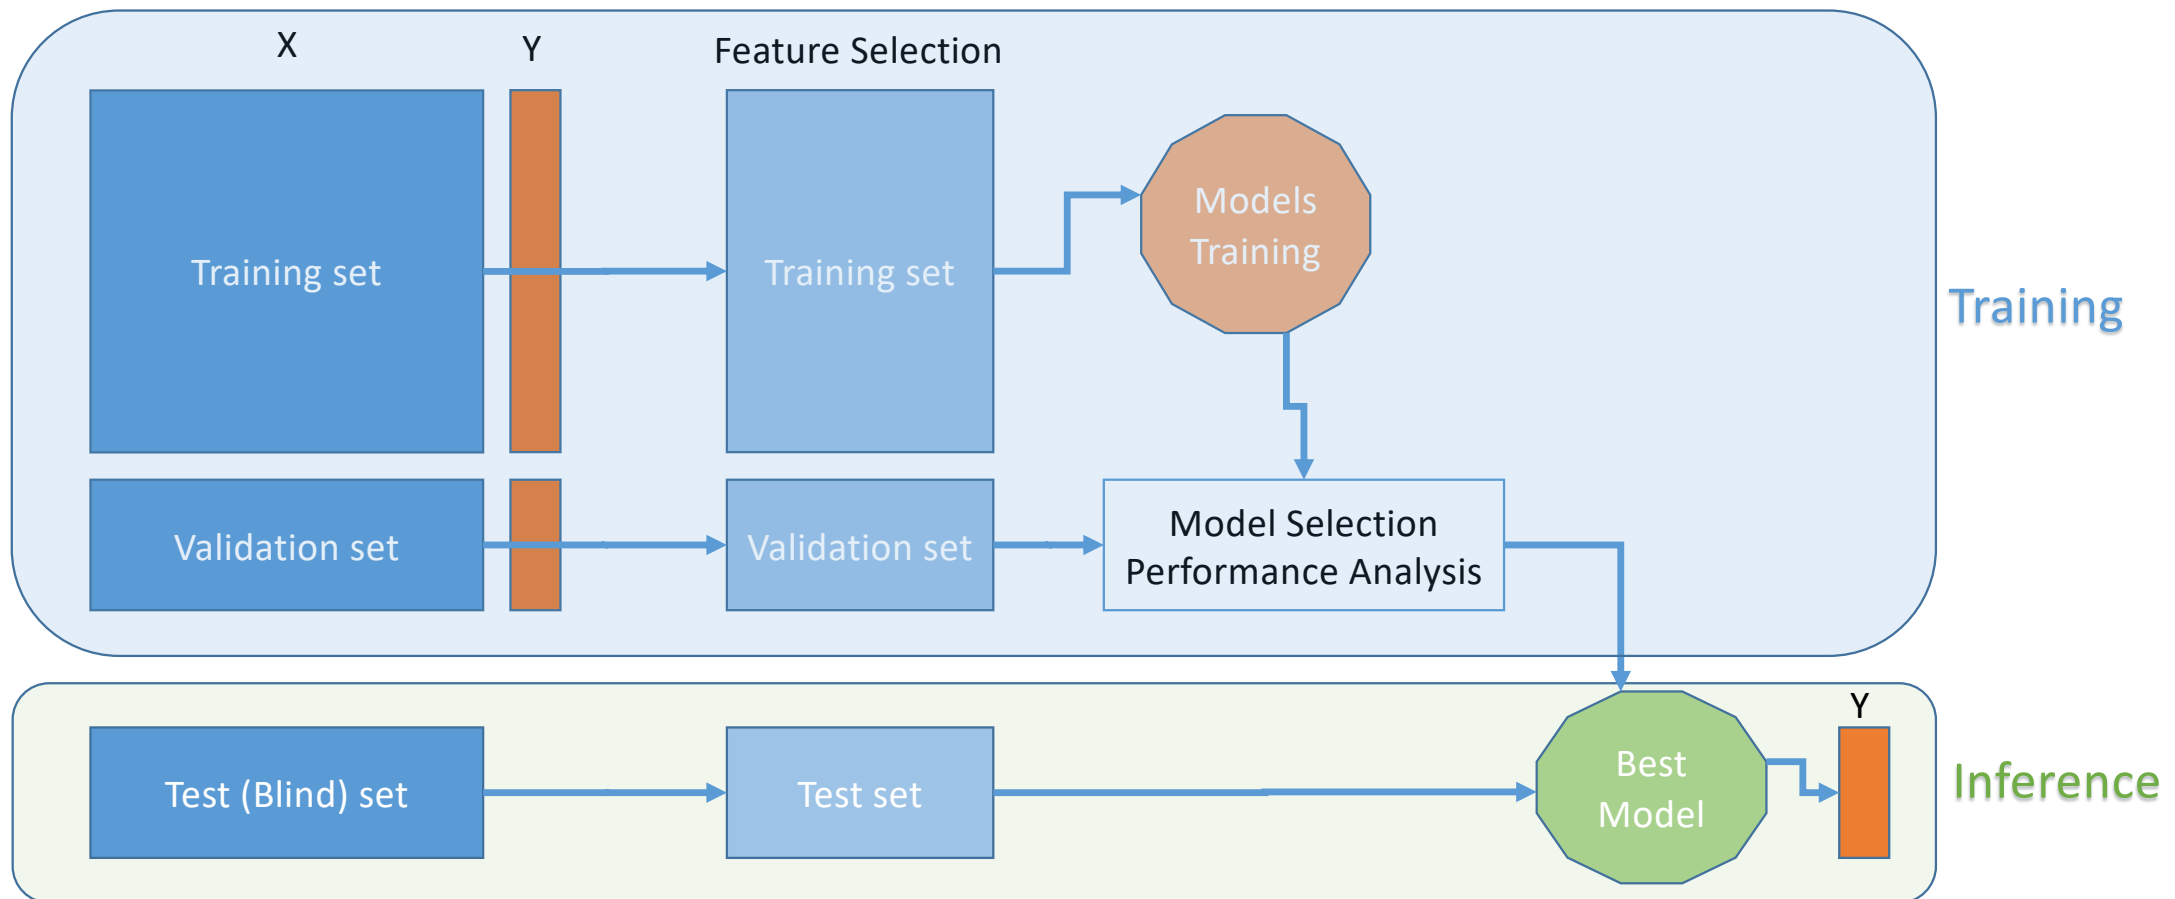
Clustering



The Big Picture!



The Big Picture!



Terminologies

- Artificial Intelligence
- Machine Learning, Deep Learning
- Statistical Learning
- Data Mining
- Deep Learning

Artificial Intelligence (1943)

- “The first work that is now generally recognized as AI was [McCullouch](#) and [Pitts](#)' 1943 formal design for [Turing-complete](#) "artificial neurons". Wikipedia
- Intelligent Machines mimics Natural Intelligence (NI)
- Natural Intelligence (General Intelligence)
 - Reasoning, Problem solving,
 - Knowledge representation, Learning,
 - Planning, Perception, Motion and manipulation, Natural Language
 - Etc.

Machine Learning (1959)

- “[Arthur Samuel](#), an American pioneer in the field of [computer gaming](#) and [artificial intelligence](#), coined the term "Machine Learning" in 1959 while at [IBM](#)”. Wikipedia
- A subfield of **Computer Science** and **Artificial Intelligence** which deals with building systems that can **learn from data**, instead of explicitly programmed instructions.
- Artificial Neural Networks (**1975**)
 - Begin in 1943, stagnated in 1969, relaunched in 1975 by the Backpropagation algorithm,
- Book: “Machine Learning”. Tom M. Mitchell. 1997

Statistical Learning (1968)

- VC Theory. “On the Uniform Convergence of Relative Frequencies of Events to Their Probabilities”. Vapnik, V. N.; Chervonenkis, A. Ya, 1968
- A subfield of **Mathematics** which deals with **finding relationship between variables** to predict an outcome
- Support Vector Machines (**1995**)
 - Much simpler, overtook ANN, Vapnik V. N.
- Book
 - “An introduction to statistical learning with applications in R” (1st Edition 2013). Gareth James, Daniela Witten, Trevor Hastie, Robert Tibshirani.


Data Mining (1990)

Appeared in the **database and financial** community to recognize customer and products trends


Definition : “The process of automatically discovering useful information in large repositories”.

- **Automatically**
 - Stats: correlation between 2 variables, what is the problem?
 - DM: parallel correlation between 1000 variables, send and email if two variables are correlated in some way.
- **Discovering useful information**
 - Stats: answer a specific question
 - DM: look for any specific reason
- **Large Repositories**
 - Stats: Collect data to answer a specific question
 - DM: Collect all, you don't know the reason yet!

Book: Introduction to Data Mining (2nd edition 2018, 1st Edition in 2005).
Pang-Ning Tan, Michael Steinbach, Anuj Karpatne, Vipin Kumar



Deep Learning



How can I Learn?

- Math
 - Statistics, Probabilistic Graphical Models, Algebra, Optimization
- Programming Languages
 - Python, R, **Julia !**
- Books
 - Gilbert Strang, Linear Algebra and Learning from Data. 2018.
 - Ian Goodfellow et al. “Deep Learning”. 2016
 - Aurélien Géron. “Hands on ML with sklearn”. 2017
 - Gareth James et al., “An introduction to statistical learning with R”. 2013
 - Etc.

How can I Learn?

- MOOCs
 - Coursera.org, Udemy.com, ocw.mit.edu, etc.
- StackOverflow
- Research Papers
 - Read and rewrite algorithms from scratch
- Follow People:
 - Androw Ng, Yann LeCun, Jeff Hinton, Sebastian Thrun, Yoshua Bengio, etc.

How can I Apply?

- Start small projects and use Frameworks
 - Scikit-learn, TensorFlow, Keras, Pytorch, Caffe, Microsoft Cognitive Toolkit (CNTK), MXNet, Spark MLlib, etc.
- Challenge your self
 - Find data: Web, UCI Machine Learning Repo
 - Go for competitions: Kaggle, DrivenData, Zindi
- Github
 - Find codes
 - Share your code
- Softwares (for non-pro !)
 - Knime, IBM SPSS Modeler